

FAP: A Foveation-Inspired Adversarial Purification Pipeline for Enhancing Robustness in Mammography Classification

Ghazal Lalooha, Wenjie Ruan, Venus Haghighi, Xinshu Li, and Quan Z. Sheng

School of Computing, Macquarie University, Sydney, NSW 2109, Australia

Email: {ghazal.lalooha, w.ruan, venus.haghighi, xinshu.li, michael.sheng}@mq.edu.au

Abstract—Deep learning models for medical image analysis demonstrate remarkable diagnostic accuracy but remain highly vulnerable to adversarial perturbations. To address this challenge, we introduce Foveated Adversarial Purification (FAP), a biologically inspired preprocessing pipeline that integrates three core innovations. First, FAP employs *eccentricity-adaptive separable Gaussian blurring*, where kernel size dynamically adjusts with lesion morphology. This approach mimics the human fovea’s acuity gradient, preserves high-frequency details around lesions while suppressing peripheral noise, and reduces GPU memory usage by 40% compared to conventional 2D filtering. Second, FAP introduces *gradient-guided fixation sampling with sigmoid-clustered probability*, which prioritizes lesion-dense regions consistent with radiologists’ diagnostic scanpaths. This mechanism achieves 82% overlap with radiologist-annotated regions of interest, ensuring that preprocessing aligns with clinical saliency rather than arbitrary regions. Third, FAP implements *lesion-aware adversarial training*, where binary spatial masks confine perturbations to non-diagnostic regions. This preserves lesion fidelity while hardening the classifier against attacks, yielding a certified ℓ_2 radius of 1.12, exceeding prior defenses. Evaluated across three mammography datasets, FAP achieves substantial robustness improvements: +20.03% absolute accuracy on CMMD (coarse tumors), +16.39% on BREAST (mixed lesions), and maintains baseline performance on CBIS-DDSM (microcalcifications). By aligning computational robustness with biological vision strategies, FAP establishes a clinically interpretable and computationally efficient framework for adversarial defense in medical imaging. The implementation is released in our GitHub repository¹.

Index Terms—Adversarial training, Medical imaging, Foveated preprocessing, Gradient-guided fixation, Clinical robustness

I. INTRODUCTION

Deep learning has revolutionized medical image analysis, offering unprecedented capabilities for disease detection, diagnosis, and treatment planning. However, its widespread clinical adoption remains significantly hindered by a critical vulnerability: susceptibility to adversarial perturbations. These perturbations are subtle, often imperceptible input manipulations designed to exploit model sensitivities to non-diagnostic high-frequency features, thereby undermining the reliability of deep learning systems in real-world clinical settings [1]–[3]. While adversarial defense research has made remarkable progress in natural image domains [4]–[6], medical imaging

introduces unique challenges that demand novel solutions [7], [8].

Specifically, three domain-specific research challenges complicate the development of effective adversarial defenses in clinical settings [9], [10]. Firstly, medical images often contain lesions with highly variable shapes, sizes, and spatial scales, spanning several orders of magnitude [11]–[13]. This variability requires defense mechanisms that can adapt to the diverse spatial characteristics of medical data [14]. Secondly, unlike natural image datasets, which are abundant and richly labeled, medical imaging datasets are typically smaller and sparsely annotated [15]–[17]. This issue restricts the applicability of conventional adversarial training techniques, which require large volumes of labeled data in order to generalize effectively [17]. Lastly, in medical imaging, it is essential that defense mechanisms do not compromise the integrity of diagnostically relevant regions [18]. Any preprocessing or adversarial mitigation strategy must ensure that these critical areas remain intact to maintain diagnostic accuracy [19].

Current adversarial defense methods, despite their promise, still fall short of addressing the aforementioned challenges comprehensively [20], [21]. Current defense strategies that rely on uniform preprocessing or global adversarial training often adopt perturbation-agnostic approaches, which fail to account for the unique characteristics of medical imaging [22], [23]. Moreover, they disregard the perceptual strategies employed by radiologists during image interpretation. Radiologists utilize a foveated visual processing approach, in which central high-acuity vision is used to analyze diagnostically relevant lesions, while peripheral vision suppresses noise through frequency-dependent filtering [24], [25]. This perceptual strategy highlights a fundamental disconnect between existing defense mechanisms and the biological principles that govern human expertise in medical image analysis.

To address these limitations, we propose a novel framework called *Foveated Adversarial Purification (FAP)*, which bridges computational robustness with insights from perceptual neuroscience. The FAP framework integrates three key components to bridge computational robustness with perceptual neuroscience:

- *Eccentricity-adaptive separable blurring* replaces isotropic filtering with spatially variant kernels that decompose 2D operations into cascaded 1D convolutions.

¹<https://github.com/ghazallalooha/FAP>

This biologically-inspired approach mimics the human visual system's acuity gradient, preserving diagnostic features in fixation zones while attenuating high-frequency artifacts in peripheral regions.

- *Gradient-guided fixation clustering* emulates radiologists' lesion prioritization through saliency-driven region sampling, generating anatomically constrained adversarial examples using hybrid loss optimization. The fixation probability at pixel (x,y) follows a sigmoid distribution, prioritizing regions with high gradient magnitudes that correlate with diagnostically salient features.
- *Lesion-focused adversarial training* employs gradient masking to confine perturbations to non-diagnostic regions, combining R-Blur [26] augmentation. The masking mechanism isolates lesion regions via binary spatial filtering.

By aligning computational robustness with perceptual neuroscience principles, the proposed FAP framework represents a novel approach to adversarial defense in medical imaging. In summary, this work contributes four advances to trustworthy AI in clinical settings:

- 1) **Biological Alignment:** We present the first integration of foveation principles into mammography AI by replacing uniform preprocessing with eccentricity-dependent kernels tuned to lesion morphology.
- 2) **Domain-Specific Training:** We design a lesion-aware adversarial training protocol that prioritizes diagnostically irrelevant regions, thereby preserving ROI fidelity while improving robustness.
- 3) **Computational Efficiency:** We reduce GPU memory consumption by 40% compared to conventional approaches through the separable blur implementation, which enables real-time deployment.
- 4) **Cross-Dataset Analysis:** We conduct a comprehensive evaluation across heterogeneous mammography datasets that identifies modality-specific limitations and informs future clinical translation.

The remainder of this paper is organized as follows. Section II reviews the relevant research on adversarial defenses in medical imaging. Section III outlines the problem and associated challenges in mammography classification. Section IV introduces the proposed Foveated Adversarial Purification (FAP) pipeline. Section V describes the experimental setup, reports the evaluation results, and conducts the results analysis. Finally, Section VI concludes the paper and discusses future directions.

II. RELATED WORK

A. Adversarial Defenses in Medical Imaging

Current medical AI defenses primarily adapt computer vision techniques, facing fundamental limitations in clinical contexts. Adversarial training, while improving robustness, degrades diagnostic accuracy on clean mammograms due to overfitting to perturbation patterns [27]–[30]. Input-space purification methods like Gaussian blurring suppress adversarial

TABLE I: Domain-Specific Defenses.

Domain	Key Mechanism	Robustness Metric
Time Series	Wavelet purification Low-freq preservation	Certified ℓ_2 radius
Finance	Transaction filtering Defense-aware ensembling	ROC AUC Harm. mean
Radiology	Lesion-aware masking Adaptive blur ($\sigma = 18.2$ mm)	PGD resistance SNR loss

noise but erase critical high-frequency features such as microcalcifications [31], [32]. Certified robustness via randomized smoothing offers theoretical guarantees [26], but becomes computationally intractable for high-resolution mammograms ($> 5,120 \times 5,120$ pixels) due to quadratic memory scaling [33]. These approaches struggle to reconcile medical imaging's dual requirements: preserving diagnostically salient features while suppressing adversarial artifacts [34].

B. Biologically-Inspired Defenses

Biological vision mechanisms have inspired robustness strategies in other domains, though medical adaptations remain nascent. R-Blur [26] simulates peripheral vision's low-frequency bias through foveated preprocessing, improving natural image robustness but requiring lesion-specific tuning for mammography. FeMPure [35], [36] demonstrates feature-space purification viability in time-series classification via wavelet-based intermediate filtering, suggesting potential for medical applications. Financial transaction models employ scanpath simulation to prioritize salient regions [37], a strategy underexplored in medical imaging despite radiologists' fixation patterns being well-characterized [38]. While effective in their native domains, these methods lack mechanisms for lesion-scale adaptability and clinical interpretability, which are critical challenges for AI applications in medical imaging.

C. Domain-Specific Robustness Strategies

Our *Foveated Adversarial Purification (FAP)* synthesizes three innovations: lesion-aware foveation with eccentricity-adaptive kernels [38], adversarial training constrained to diagnostically irrelevant regions [39], and pathologically grounded dataset selection mirroring financial attack-defender dynamics [37], [40]. As shown in Table I, lesion-aware masking with adaptive blurring functioning as the key mechanism, outperforms the previous related works.

D. Key Differentiators

Current approaches such as input-space purification and certified robustness via randomized smoothing—either sacrifice lesion fidelity or become computationally intractable for high-resolution mammograms ($> 5,120 \times 5,120$ pixels). To address this gap, Foveated Adversarial Purification (FAP) advances mammography classification robustness by integrating three core innovations. First, it implements lesion-aware adversarial training by using gradient-guided perturbation masking to restrict adversarial attacks (via Projected Gradient Descent)

to non-lesion regions, thereby protecting diagnostically critical areas while still hardening the model against perturbations—this leads to substantial accuracy improvements (over 20% and 16% absolute gains on specific datasets) compared to baselines and outperforms prior biological defenses. Second, FAP introduces gradient-guided fixation sampling, which mimics radiologists’ diagnostic scanpaths by focusing pre-processing on lesion-dense regions through sigmoid-clustered sampling, in contrast to fixed Gaussian blurring methods that lack anatomical precision. Third, FAP employs morphology-adaptive kernel sizing, dynamically adjusting convolution kernels based on lesion size and distance from fixation points. This approach reduces GPU memory usage by 40% while maintaining robustness and preserving fine image details relevant to diagnosis. Together, these innovations ensure both computational efficiency and strong defense against adversarial attacks, all while aligning closely with clinical diagnostic priorities.

III. PROBLEM STATEMENT

Deep learning models have revolutionized mammography classification. Yet, their clinical deployment remains hindered by vulnerability to adversarial perturbations, subtle, engineered distortions that exploit model sensitivities to high-frequency features absent in human diagnostic reasoning. Conventional adversarial defenses, adapted from natural imaging domains, face three critical limitations in medical contexts including Lesion Heterogeneity, Data Scarcity, and Clinical Misalignment. Fixed preprocessing techniques (e.g., uniform Gaussian blur) inadequately handle the mammographic findings (i.e., images) with very different sizes, from coarse tumors (CMMD: 18.2 mm) to microcalcifications (CBIS-DDSM: 0.8 mm), erasing diagnostically critical high-frequency features. Annotated datasets (e.g., CMMD: 1,200 cases, BREAST: 2,500 cases) are orders of magnitude smaller than natural image corpora, limiting the effectiveness of standard adversarial training. Existing defenses degrade diagnostic accuracy by perturbing lesion-dense regions and lack interpretability, failing to emulate radiologists’ biological robustness mechanisms like foveated sampling and saccadic prioritization.

IV. FOVEATED ADVERSARIAL PURIFICATION

We propose *Foveated Adversarial Purification* (FAP), a preprocessing pipeline that integrates lesion-aware adversarial training, gradient-guided fixation sampling, and morphology-adaptive kernel sizing (Table II). The pipeline is designed to address lesion heterogeneity and clinical alignment in mammography classification, providing a biologically inspired foundation for subsequent robustness evaluation (see Sec. V).

As shown in Fig. 1, the FAP pipeline operates through four stages. First, fixation sampling prioritizes lesion-dense regions using gradient-guided saliency maps, mimicking radiologists’ diagnostic scanpaths. Second, adaptive blurring applies separable Gaussian kernels dynamically tuned to tumor size distributions, preserving diagnostically critical features




Fig. 1: **FAP Conceptual Framework.** A central core of three parallel modules, fixation sampling, eccentricity-adaptive blurring, and lesion-aware masking, supports two contexts: (b) masked-PGD training to obtain robust weights and (c) inference to produce risk and CAMs.

while suppressing peripheral noise. Third, lesion-aware adversarial training confines perturbations to non-lesion regions via gradient masking, thereby protecting clinically relevant areas. Finally, a VGG16 classifier [27] is trained with a hybrid loss function to integrate both clean and adversarial examples.

A. Lesion-Aware Adversarial Training

Standard adversarial training often degrades diagnostically critical regions. To address this, FAP confines perturbations via gradient-guided masking:

$$M(x, y) = \begin{cases} 0 & \text{if } (x, y) \in \text{ROI,} \\ 1 & \text{otherwise} \end{cases} \quad (1)$$

This binary mask isolates lesion regions (ROIs) identified via class activation maps (CAMs) or gradient saliency during training. By setting masked regions to 0, adversarial perturbations are confined to non-diagnostic areas, thereby preserving lesion fidelity and aligning with radiologists’ diagnostic prioritization of lesion-dense regions.

Conventional adversarial training applies perturbations uniformly across the image space, which inadvertently degrades lesion-rich areas and reduces diagnostic accuracy. Lesion-aware masking overcomes this limitation by confining perturbations to non-ROI regions while shielding diagnostically salient structures. This ensures that adversarial examples do not distort subtle features essential for cancer detection, preserving diagnostic safety.

Projected Gradient Descent (PGD) attacks are then restricted to masked regions:

$$\delta_{t+1} = \text{Proj}_\epsilon(\delta_t + \alpha \cdot \text{sign}(\nabla_x \mathcal{L}) \odot M) \quad (2)$$

$$\epsilon = 0.004, \quad \alpha = 0.001, \quad t_{\max} = 10 \quad (3)$$

TABLE II: **Comparison of the FAP Components.** Each component is described in terms of its objective, implementation, and empirical validation.

Component	Objective	Implementation	Empirical Validation
Lesion-Aware Adversarial Training	Confine perturbations to non-diagnostic regions.	Binary masking and PGD attack restriction.	+20.03% accuracy gain on coarse tumors.
Gradient-Guided Fixation Sampling	Prioritize lesion-dense regions.	Sigmoid probability based on gradient magnitudes.	82% spatial overlap with radiologist ROIs.
Morphology-Adaptive Kernel Sizing	Adjust blur strength to lesion morphology.	Adaptive standard deviation based on lesion diameter.	+6.85% accuracy gain from separable blurring.

PGD perturbations are applied only to non-lesion regions via element-wise multiplication with M . This design ensures that adversarial noise targets anatomically irrelevant areas while shielding diagnostically important pixels. The empirical effectiveness of this strategy is demonstrated in the experimental results (see Sec. V).

B. Gradient-Guided Fixation Sampling

Conventional random or uniformly distributed fixation points fail to capture the structured scanpaths radiologists employ when interpreting mammograms. Radiologists naturally prioritize regions of high diagnostic saliency, such as masses, distortions, or calcification clusters, while deprioritizing homogeneous tissue.

To emulate this behavior, FAP replaces random fixation with probabilistic sampling based on gradient magnitudes:

$$P(x, y) = \frac{1}{1 + e^{-\gamma \cdot \|\nabla_x \mathcal{L}(x, y)\|}}, \quad \gamma = 0.1 \quad (4)$$

This sigmoid-based function modulates fixation probability such that areas with stronger gradients, which correlate with diagnostically salient features, are more likely to be selected. In this way, fixation sampling operationalizes radiologists' diagnostic prioritization by directing attention toward lesion-dense regions.

Random or uniformly distributed fixation points fail to capture the structured scanpaths radiologists naturally employ. Radiologists focus disproportionately on regions of high diagnostic saliency, while deprioritizing homogeneous tissue. Gradient-guided fixation sampling embeds this principle by assigning higher fixation probabilities to areas with strong gradient magnitudes, thereby aligning adversarial mitigation with clinically meaningful priorities.

The empirical evaluation of this strategy, including its effectiveness in aligning with radiologist annotations, is presented in Sec. V.

C. Morphology-Adaptive Kernel Sizing

Lesions in mammography vary substantially in size and morphology, ranging from coarse tumors to fine microcalcifications. Uniform Gaussian blurring indiscriminately attenuates high-frequency details, which can lead to the loss of diagnostically critical features.

To address this, FAP introduces morphology-adaptive kernel sizing, where the blur strength adapts smoothly to lesion characteristics:

$$\sigma(d) = \beta \cdot \text{sigmoid}(\alpha \cdot (d - \mu_{\text{lesion}})) \quad (5)$$

where d is the Euclidean distance from fixation points and μ_{lesion} is a dataset-specific lesion diameter parameter. The parameters α and β control the steepness and scale of the kernel transition.

This design ensures that high-frequency details are preserved near lesion regions while peripheral areas are smoothed, mimicking the acuity gradient of the human visual system. The empirical evaluation of this component, including its computational benefits, is reported in Sec. V.

Uniform Gaussian blurring indiscriminately attenuates high-frequency details, which can result in the loss of diagnostically critical structures, such as microcalcifications or fine spiculations. Morphology-adaptive kernel sizing addresses this by varying kernel size as a function of distance from fixation points, mirroring the fovea's acuity gradient. This balances robustness against adversarial perturbations with preservation of diagnostic fidelity.

D. Clinical Integration

A key objective of FAP is to ensure seamless integration into existing clinical workflows. In a typical deployment, digital mammograms are first retrieved from the Picture Archiving and Communication System (PACS). FAP preprocessing is applied immediately after retrieval, so that images are purified from potential adversarial perturbations before being passed to the classification model. The classifier then produces a malignancy risk score and a heatmap localization, which are forwarded to the radiologist for interpretation.

This workflow preserves the radiologist's central role in decision-making while adding an additional layer of automated robustness. The practical implications of integrating FAP into clinical pipelines, including efficiency and case-based analyses, are further explored in Sec. V.

Together, these methodological design choices are necessary to reconcile the trade-off between robustness and diagnostic integrity, ensuring that adversarial defense in mammography remains biologically plausible, computationally efficient, and clinically aligned.

1) Clinical Relevance and Case Study: The ultimate goal of developing robust mammography classifiers is not only to improve benchmark robustness but also to support radiologists in real-world screening workflows. To illustrate the practical

impact of FAP, we consider a case study that reflects a typical deployment pipeline in a hospital setting.

2) *Integration into Clinical Workflow*: Digital mammograms are first retrieved from the Picture Archiving and Communication System (PACS). FAP preprocessing is applied immediately after retrieval, ensuring that incoming images are purified from potential adversarial perturbations before being passed to the classification model. The classifier then produces a malignancy risk score and heatmap localization, which are forwarded to the radiologist for interpretation. This workflow preserves the radiologist’s central role in decision-making while adding an additional layer of automated robustness.

An ablation of these components is presented in Sec. V-F and a qualitative preview of their visual effect appears in Fig. 4.

V. EXPERIMENTAL STUDIES

A. Datasets and Preprocessing

We evaluate the FAP framework on three publicly available mammography datasets from The Cancer Imaging Archive (TCIA)², each selected to represent distinct lesion morphologies and imaging modalities.

CMMD (Chinese Mammography Database): Comprising 1,200 digital mammography images with histopathology-confirmed labels, CMMD focuses on coarse tumor detection (median lesion size: 18.2 mm). Images are converted from DICOM to PNG format and processed using stratified splitting (60/20/20). During preprocessing, eccentricity-adaptive blurring is applied as defined in Eq. (5), and fixation sampling follows Eq. (4), achieving 82% spatial overlap with radiologist-annotated ROIs.

BREAST Screening DBT dataset: Includes 2,500 cases of mixed lesions (masses, distortions) from digital breast tomosynthesis (DBT) and synthesized 2D projections. Confirmed malignant cases undergo histogram equalization, followed by a 70/15/15 split. FAP’s separable blur implementation (Eq. (5)) reduces GPU memory consumption by 40%, critical for high-resolution DBT volumes. Lesion-aware adversarial training is applied, confining PGD perturbations to non-ROI regions (Eqs. (1), (2)–(3)), contributing to a 16.39% absolute accuracy gain over baseline models.

CBIS-DDSM: Contains 3,000 digitized film mammograms with microcalcification annotations (median size: 0.8 mm). ROI cropping and wavelet denoising are applied to isolate sub-1.2 mm microcalcifications, though irreversible SNR loss occurs due to the kernel’s FWHM exceeding lesion dimensions. Despite this limitation, FAP maintains baseline accuracy, underscoring the need for future high-frequency adaptive extensions.

All datasets undergo the FAP-specific preprocessing steps summarized in Table II: eccentricity-adaptive blurring (Eq. (5)), gradient-guided fixation sampling (Eq. (4)), and lesion-aware masking during adversarial training (Eqs. (1), (2)–(3)). Computational efficiency metrics (15 ms/image on

RTX 4090 GPUs) ensure clinical deployability while preserving patient confidentiality.

B. Implementation Details

The FAP pipeline employs separable Gaussian blur with horizontal/vertical kernel decomposition to handle high-resolution mammograms ($> 5,120 \times 5,120$ pixels). The kernel standard deviation adapts dynamically to lesion morphology, as defined in Eq. (5) (see Sec. IV-C). This adaptation reduces GPU memory consumption by 40% compared to conventional 2D convolution.

Fixation points follow the sigmoid-based probability distribution introduced in Eq. (4) (Sec. IV-B), achieving 82% spatial overlap with radiologist annotations (Fig. 2). While the CMMD dataset implements random sampling, the experimental results confirm the efficacy of gradient-guided prioritization.

Lesion-aware adversarial training confines PGD attacks to non-lesion regions via binary masking (Eq. (1)) and the PGD update rule with parameters (Eqs. (2)–(3), Sec. IV-A). This design contributes to a 16.39% accuracy gain on BREAST under PGD attacks.

We train the models using the Adam optimizer [41] with a learning rate of 10^{-4} and a weight decay coefficient of 10^{-5} , ensuring stable gradient updates while mitigating overfitting in data-scarce medical contexts. This configuration enables effective adaptation to mammographic feature complexity without parameter drift, as validated by $< 1\%$ training loss variance across epochs in ablation studies.

We combine clean and adversarial example contributions using the following hybrid loss function:

$$L = 0.7L_{\text{clean}} + 0.3L_{\text{adv}}, \quad (6)$$

where L_{clean} and L_{adv} denote cross-entropy losses for unperturbed and adversarially modified mammograms, respectively. This weighting preserves diagnostic fidelity (78.77% baseline accuracy on CMMD) while hardening against ℓ_{∞} -bounded perturbations.

Data augmentation includes random horizontal/vertical flips ($p = 0.5$) and R-Blur [26] transformations with three fixation points per image. The flip probability induces anatomical variability across eight possible orientations, approximating clinical acquisition diversity. FAP emulates radiologists’ foveated processing through multi-focal perspective generation (3 fixation points/image), eccentricity-dependent blurring (Eq. (5)), and gradient-guided lesion prioritization (Eq. (4)).

This combined strategy expands effective training data $8\times$ while suppressing high-frequency adversarial artifacts through biologically inspired filtering. The modular implementation achieves 15 ms/image throughput on RTX 4090 GPUs, ensuring clinical workflow compatibility.

All experimental evaluations are conducted using NVIDIA GeForce RTX 4090 GPUs with 24 GB of GDDR6X memory, which provides the computational resources necessary for training and inference on high-resolution mammography images. The FAP pipeline is implemented in PyTorch 2.0,

²<https://www.cancerimagingarchive.net/>

TABLE III: Overview of Mammography Datasets, Lesion Characteristics, and Preprocessing Strategies

Dataset	Description	Data Used in Our Study	Preprocessing & Splitting
CMMD	Digital 2D Pathology Coarse tumors 1,200 cases	All images	DICOM→PNG Stratified split (60/20/20) All images processed
BREAST Screening DBT	DBT & synth. 2D Pathology + radiologist Mixed lesions 2,500 cases	2D projections, confirmed	Histogram equalization Confirmed cases 70/15/15 split
CBIS-DDSM	Digitized film 2D Pathology + ROI masks Microcalcifications 3,000 cases	ROI cropping, confirmed	ROI cropping Wavelet denoising Confirmed cases 65/15/20 split

TABLE IV: Robustness Gains Across Datasets.

Dataset (Lesion Type)	FAP Δ	Baseline Δ
CMMD (Coarse tumors)	+20.03%	+3.38% (R-Blur)
BREAST (Mixed lesions)	+16.39%	+6.07% (AT)
CBIS-DDSM (Microcalc.)	+0.2%	Neutral




Fig. 2: Visualization of scanpath fixation points on a mammogram. Red markers and connecting lines indicate the sequential locations and the order of model-predicted visual fixations, demonstrating the clustering of attention around diagnostically salient regions. This scanpath emulates radiologists’ diagnostic strategies, highlighting FAP’s alignment with clinical perception.

leveraging CUDA 12.1 for accelerated deep learning operations, ensuring compatibility with the latest GPU architectures and efficient utilization of tensor cores for both training and real-time inference. This hardware configuration supports the modular design of the pipeline, allowing for direct comparison between biological and non-biological defenses while maintaining diagnostic fidelity across heterogeneous lesion scales.

C. Adversarial Threat Models

We evaluate the adversarial robustness of the Foveated Adversarial Purification (FAP) framework under white-box and black-box attack scenarios, simulating real-world clinical threat models. Three attack types are implemented with lesion-aware constraints, as shown in Table V.

Projected Gradient Descent (PGD): We employ ℓ_∞ -bounded perturbations with the update rule and parameter settings defined in Eqs. (2)–(3). Gradient-guided masking ensures that perturbations are confined to non-lesion regions, where M isolates irrelevant pixels using class activation maps (CAMs). This achieves 82% spatial overlap with radiologist annotations.

Training Integration: Adversarial examples are generated on-the-fly using the hybrid loss defined in Eq. (6). This ensures that both clean and adversarial contributions are balanced during training, preserving diagnostic fidelity while improving robustness.

Certified Robustness: We compute the certified ℓ_2 radius with 99% confidence using randomized smoothing:

$$R = \sigma \cdot \Phi^{-1}(p_A). \quad (7)$$

FAP achieves a certified radius of 1.12, which is 12% larger than R-Blur (1.00) and 8% larger than adversarial training (1.04), demonstrating stronger theoretical guarantees against adversarial perturbations.

The adversarial perturbation bounds (Eq. (3)) are anatomically aligned with lesion size distributions across datasets (18.2 mm for CMMD and 12.7 mm for BREAST), ensuring clinically plausible noise injection while preserving diagnostically critical features. Gradient-guided masking (Eq. (1)) confines perturbations to non-lesion regions, achieving 82% spatial overlap with radiologist-annotated ROIs. Computational efficiency is maintained through separable Gaussian blur decomposition (Eq. (5)), reducing GPU memory consumption by 40% while enabling real-time operation (15 ms/image) on RTX 4090 GPUs. Together, these design choices ensure robustness against ℓ_∞ -bounded attacks while maintaining diagnostic fidelity across heterogeneous lesion scales.

TABLE V: White-Box Attack Parameters.

Attack (Norm)	ϵ	Steps	Adaptation
PGD (ℓ_∞)	0.004	10	Lesion masking
FGSM (ℓ_∞)	0.008	1	Grad. clip ($\pm 3\sigma$)
CW (ℓ_2)	1.0	100	Conf. reduction ($\kappa=0$)




Fig. 3: **Ablation on model components.** Accuracy on CMMD and BREAST datasets for the baseline model and after adding separable blur, fixation sampling, and lesion-aware training.

D. Baselines

Baseline Architecture: We adopt VGG16 [27] as the baseline classifier due to its well-established performance in medical image analysis and its suitability for high-resolution mammograms. The network’s sequential convolutional layers provide strong feature extraction capability while maintaining interpretability through class activation maps (CAMs). Using VGG16 [27] also ensures comparability with prior adversarial robustness studies in medical imaging, making it an appropriate reference point for evaluating FAP’s improvements. We compare FAP against both biological and non-biological defenses as baselines. Biological baselines feature R-Blur [26] with a fixed Gaussian blur ($\sigma = 0.12$) and R-Warp [42], which applies foveated resampling with DCT compression. Non-biological baselines comprise adversarial training (AT) [43] using PGD ($\epsilon = 0.004$, 7 steps) and certified robustness via randomized smoothing ($\sigma = 0.125$) [44].

E. Evaluation Metrics

We assess model performance using three rigorously defined metrics aligned with clinical and computational robustness requirements.

Clean accuracy quantifies baseline diagnostic fidelity on unperturbed test sets, measured via ROC AUC. Robust accuracy evaluates resilience to adversarial perturbations, computed as the fraction of test samples correctly classified under ℓ_∞ -bounded attacks (e.g., iterative PGD with $\epsilon = 0.004$). Certified robustness is assessed through randomized smoothing, which estimates the maximum ℓ_2 perturbation radius within

TABLE VI: Cross-Dataset Robustness Comparison (%).

Method	CMMD	BREAST	CBIS-DDSM
Baseline (VGG16 [27])	78.77	72.34	81.09
R-Blur [26]	82.15	75.62	80.95
AT [43]	85.23	78.41	82.67
FAP (Ours)	96.80	88.73	81.24

which predictions remain provably stable at 99% confidence, following the Clopper-Pearson methodology.

Together, these metrics provide a comprehensive assessment of diagnostic fidelity, adversarial robustness, and provable safety guarantees, forming the basis for the quantitative analysis presented in Section V-F.

F. Results

Before turning to quantitative results, Fig. 4 visualizes the effect of the proposed preprocessing on representative cases from CMMD (coarse tumors), BREAST (mixed lesions), and CBIS-DDSM (microcalcifications). The images illustrate how FAP suppresses peripheral clutter while preserving lesion-dense regions.

As suggested qualitatively in Fig. 4, datasets dominated by larger masses (CMMD, BREAST) benefit most, whereas CBIS-DDSM with sub-millimeter calcifications shows limited improvement. We present cross-dataset robustness comparisons in Table VI. From the table, we can see that FAP achieves the state-of-the-art performance on CMMD and BREAST, with absolute improvements of 20.03% and 16.39% over the baseline, respectively. On CBIS-DDSM, FAP maintains baseline accuracy, reflecting its limitation on high-frequency microcalcifications.

The cross-dataset robustness improvements of FAP compared to baselines are summarized in Table IV.

An ablation study (see Table VII) demonstrates the incremental contribution of each FAP component. Separable blur, fixation sampling, and lesion-aware training collectively drive robustness gains, with the full pipeline achieving the highest accuracy on both CMMD and BREAST.

1) Case Example: Consider a screening scenario in which a 52-year-old patient presents for a routine mammogram. A baseline classifier under adversarial perturbation misclassifies a malignant mass as benign due to subtle noise injected into the peripheral tissue. By contrast, FAP’s eccentricity-adaptive blurring preserves high-resolution detail around the lesion, while gradient-guided fixation directs attention to the region of interest. The lesion-aware mask further ensures that adversarial noise cannot distort diagnostically critical pixels. The resulting risk score correctly flags the abnormality, enabling timely referral for biopsy.

This case highlights three clinically relevant advantages of FAP: (1) diagnostic safety is maintained by preventing distortion of lesion-rich areas, (2) efficiency is achieved since preprocessing is lightweight and can be integrated into PACS-based workflows, and (3) interpretability is enhanced, as fixation maps can be visualized alongside heatmaps to reassure




Fig. 4: Qualitative effect of the proposed Foveated Adversarial Purification (FAP) on representative mammograms from the three datasets used in this work. Each pair is arranged with the original image on the **left** and the corresponding preprocessed image on the **right**. From top to bottom: (i) **CMMD (coarse tumors)**: peripheral clutter is suppressed while lesion morphology and margins are preserved. (ii) **BREAST (mixed lesions)**: background texture is attenuated and lesion-dense regions retain contrast. (iii) **CBIS-DDSM (microcalcifications)**: high-frequency structures are largely preserved, although peripheral smoothing may diminish very fine detail, consistent with the more limited robustness gains observed on this dataset.

radiologists that the model focuses on clinically meaningful regions. These properties suggest that FAP is not merely a technical improvement but a step toward trustworthy deployment of adversarially robust AI in mammography screening.

2) *Certified Robustness*: We show that FAP’s certified ℓ_2 radius significantly exceeds that of the baselines (Table VIII), demonstrating stronger theoretical guarantees against adversarial perturbations. FAP achieves a certified radius of 1.12, which represents a 12% improvement over both R-Blur and Adversarial Training (AT).




Fig. 5: **Clean and Robust Accuracy by Method and Dataset.**

We demonstrate that FAP not only improves empirical robustness under adaptive attack scenarios but also provides provable lower bounds on robustness through the certification process. In other words, we guarantee that any adversarial perturbation within this radius cannot alter the model’s prediction, which ensures predictable performance even under worst-case conditions.

By exceeding the certified radii of the baselines, we establish a more reliable safety margin with FAP. This result increases confidence that our purification mechanism works effectively in practice and in theory, offering robustness that extends beyond specific attack settings.

TABLE VII: Ablation Study (CMMD).

Component	Accuracy Gain
Baseline (VGG16)	78.77%
+ Separable Blur	+6.85%
+ Fixation Sampling	+12.57%
+ Lesion-Aware Training	+18.03%

3) *Failure Case Analysis*: We analyze microcalcifications in the CBIS-DDSM dataset and find that their median size of 0.8mm falls below FAP’s blur kernel full width at half maximum (FWHM) of 1.2mm. Because of this mismatch, FAP attenuates the fine-grained features that characterize these lesions and removes critical diagnostic detail.

We observe that this outcome aligns with radiologists’ reduced sensitivity to sub-kernel lesions [45]. Human experts also struggle to detect calcifications smaller than the effective blur scale, which shows that the difficulty stems from perceptual and resolution limits rather than from FAP alone.

Based on this analysis, we argue that defenses should adapt to lesion frequency and scale. Instead of applying a uniform kernel across the image, we recommend developing frequency-adaptive or multi-resolution purification strategies. Such approaches can preserve diagnostically important microcalcifications, while still suppressing adversarial perturbations, reducing the trade-off between robustness and sensitivity, and

TABLE VIII: Certified Robustness and Attack Parameters.

Method	Attack	ϵ	Steps	Certified ℓ_2
R-Blur [26]	PGD (ℓ_∞)	0.03	10	1.00
AT [43]	PGD (ℓ_2)	0.12	100	1.04
FAP (Ours)	PGD (ℓ_∞)	0.004	10	1.12

improving the clinical applicability of defense mechanisms in medical imaging.

VI. CONCLUSION AND FUTURE DIRECTIONS

This work introduces *Foveated Adversarial Purification (FAP)*, a biologically inspired preprocessing pipeline that can significantly enhance adversarial robustness in mammography classification. The framework integrates three key innovations: (1) Eccentricity-adaptive Gaussian blurring that dynamically adjusts kernel size based on tumor morphology, achieving $\Delta + 20.03\%$ accuracy gain on CMMD (coarse tumors) through the kernel parameterization, (2) Sigmoid-clustered fixation sampling that mimics radiologists' diagnostic scanpaths with 82% spatial overlap ($SS = 0.82 \pm 0.04$), prioritizing lesion-dense regions through gradient-guided probabilistic clustering, and (3) Lesion-aware adversarial training, which confines perturbations to non-ROI areas, yielding a certified ℓ_2 radius of 1.12—12% larger than prior methods. This addresses medical imaging's triad of challenges: lesion heterogeneity, data scarcity, and clinical interpretability requirements.

FAP demonstrates clinical viability with a 15 ms/image processing overhead on RTX 4090 GPUs, compatible with real-time screening workflows. The method shows strong generalization for tumors ≥ 5 mm (BREAST: $+16.39\%$ accuracy) but remains limited for sub-resolution microcalcifications < 1 mm (CBIS-DDSM: $\Delta + 0.2\%$). Unlike global purification approaches, FAP's lesion-focused masking preserves diagnostic ROI fidelity while maintaining computational tractability.

Although the proposed FAP framework demonstrates promising improvements in adversarial robustness for mammography classification, several research avenues remain open for exploration.

Multi-Modal Imaging Integration: Mammography is the primary modality for breast cancer screening, yet complementary modalities such as ultrasound and magnetic resonance imaging (MRI) are frequently employed in diagnostic follow-up. We plan to investigate extending FAP to a multi-modal setting in future work, where robustness mechanisms are adapted to modality-specific characteristics while leveraging cross-modal consistency. This direction may reduce false negatives by ensuring that adversarial perturbations targeting one modality cannot undermine the overall diagnostic pipeline.

Federated and Privacy-Preserving Learning: In real-world healthcare environments, sensitive patient data are distributed across institutions and cannot be centralized due to privacy and legal constraints. Incorporating FAP into federated learning frameworks would enable collaborative robustness training without direct data sharing. Moreover, designing privacy-preserving preprocessing steps aligned with regulations such

as HIPAA and GDPR remains an important step toward the clinical adoption of adversarially robust models.

Regulatory and Translational Challenges: Clinical deployment of AI systems requires compliance with stringent regulatory frameworks, including approval processes by agencies such as the FDA and TGA. Beyond technical robustness, FAP and related defenses must demonstrate reproducibility, reliability under distributional shifts, and transparent interpretability. Addressing these requirements will be essential for ensuring that robustness improvements translate from research prototypes to regulatory-cleared medical devices.

Dataset Shift and Generalization: A major challenge in medical imaging is the performance degradation that occurs when models are deployed on data from new scanners, hospitals, or patient populations. Future research should evaluate FAP under such domain shifts and explore adaptive mechanisms that maintain robustness in heterogeneous clinical environments. Approaches such as continual learning and domain generalization may provide a pathway to sustaining robust performance in diverse deployment scenarios.

We believe that advancing FAP along these directions will bridge the gap between methodological innovation and clinical translation, paving the way for trustworthy and scalable deployment of adversarially robust AI in healthcare.

ACKNOWLEDGMENT

This research was supported in part by the Australian Research Council (ARC) through a Discovery Early Career Researcher Award (DECRA). The work of Ghazal Laloocha was also supported in part by the International Macquarie Research Excellence Scholarship Program (iMQRES).

REFERENCES

- [1] E. Abdukhamidov, M. Abuhamad, S. S. Woo, E. Chan-Tin, and T. Abuhmed, "Hardening interpretable deep learning systems: Investigating adversarial threats and defenses," *IEEE Transactions on Dependable and Secure Computing*, vol. 21, no. 4, pp. 3963–3976, 2024.
- [2] Y. Wang, Y. Wu, S. Wu, X. Liu, W. Zhou, L. Zhu, and C. Zhang, "Boosting the transferability of adversarial attacks with frequency-aware perturbation," *IEEE Transactions on Information Forensics and Security*, vol. 19, pp. 6293–6304, 2024.
- [3] W. S. Burton II, "Applications of computer vision in biomechanics and orthopaedics," Ph.D. dissertation, University of Denver, November 2024. [Online]. Available: <https://digitalcommons.du.edu/etd/2517>
- [4] Y. Chen, X. Li, X. Wang, P. Hu, and D. Peng, "Diffilter: Defending against adversarial perturbations with diffusion filter," *IEEE Transactions on Information Forensics and Security*, vol. 19, pp. 6779–6794, 2024.
- [5] S. Qiu, Q. Liu, S. Zhou, and W. Huang, "Adversarial attack and defense technologies in natural language processing: A survey," *Neurocomputing*, vol. 492, pp. 278–307, 2022.
- [6] J. Dong, J. Chen, X. Xie, J. Lai, and H. Chen, "Survey on adversarial attack and defense for medical image analysis: Methods and challenges," *ACM Computing Surveys*, vol. 57, no. 3, pp. 1–38, 2024.
- [7] K. Vrettos, E. Koltsakis, A. H. Zibis, A. H. Karantanias, and M. E. Klontzas, "Generative adversarial networks for spine imaging: A critical review of current applications," *European Journal of Radiology*, vol. 171, p. 111313, 2024.
- [8] A. W. Salehi, S. Khan, G. Gupta, B. I. Alabdullah, A. Almjally, H. Alsolai, T. Siddiqui, and A. Mellit, "A study of cnn and transfer learning in medical imaging: Advantages, challenges, future scope," *Sustainability*, vol. 15, no. 7, p. 5930, 2023.

- [9] E. Darzi, F. Dubost, N. M. Sijtsma, and P. M. van Ooijen, "Exploring adversarial attacks in federated learning for medical imaging," *IEEE Transactions on Industrial Informatics*, vol. 20, 2024.
- [10] A. Ponsiglione and R. Cuocolo, "Radiology under siege? adversarial attacks against deep learning algorithms," *European Journal of Radiology*, vol. 169, 2023.
- [11] Y. Wang, L. Zhen, J. Zhang, M. Li, L. Zhang, Z. Wang, Y. Feng, Y. Xue, X. Wang, Z. Chen, T. Luo, R. S. M. Goh, and Y. Liu, "Mednas: Multiscale training-free neural architecture search for medical image analysis," *IEEE Transactions on Evolutionary Computation*, vol. 28, no. 3, pp. 668–681, 2024.
- [12] L. Luo, Y. Li, Z. Chai, H. Lin, P.-A. Heng, and H. Chen, "Scale-aware super-resolution network with dual affinity learning for lesion segmentation from medical images," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 36, 2024.
- [13] H. Du, J. Wang, M. Liu, Y. Wang, and E. Meijering, "Swinpa-net: Swin transformer-based multiscale feature pyramid aggregation network for medical image segmentation," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 35, no. 4, pp. 5355–5366, 2022.
- [14] X. Shen, H. Huang, B. Nichyporuk, and T. Arbel, "Improving robustness and reliability in medical image classification with latent-guided diffusion and nested-ensembles," *IEEE Transactions on Medical Imaging*, 2025.
- [15] C. Ouyang, C. Biffi, C. Chen, T. Kart, H. Qiu, and D. Rueckert, "Self-supervised learning for few-shot medical image segmentation," *IEEE Transactions on Medical Imaging*, vol. 41, no. 7, pp. 1837–1848, 2022.
- [16] X. Li, X. Ouyang, J. Zhang, Z. Ding, Y. Zhang, Z. Xue, F. Shi, and D. Shen, "Carotid vessel wall segmentation through domain aligner, topological learning, and segment anything model for sparse annotation in mr images," *IEEE Transactions on Medical Imaging*, vol. 43, no. 12, pp. 4483–4495, 2024.
- [17] M. Götz, C. Weber, F. Binczyk, J. Polańska, R. Tarnawski, B. Bobek-Billewicz, U. Köthe, J. Kleesiek, B. Stieltjes, and K. H. Maier-Hein, "Dalsa: Domain adaptation for supervised learning from sparsely annotated mr images," *IEEE Transactions on Medical Imaging*, vol. 35, no. 1, pp. 184–196, 2016.
- [18] L. Alzubaidi, A.-D. Khamael, H. A.-H. Obeed, A. Saihood, M. A. Fadhel, S. A. Jebur, Y. Chen, A. S. Albahri, J. Santamaría, A. Gupta, and Y. Gu, "Meff – a model ensemble feature fusion approach for tackling adversarial attacks in medical imaging," *Intelligent Systems With Applications*, vol. 22, p. 200355, 2024.
- [19] Y. Yang, H. Zhang, J. W. Gichoya, D. Katabi, and M. Ghassemi, "The limits of fair medical imaging ai in real-world generalization," *Nature Medicine*, vol. 30, no. 10, pp. 2838–2848, 2024.
- [20] Y. Lei, C. Niu, J. Zhang, G. Wang, and H. Shan, "Ct image denoising and deblurring with deep learning: Current status and perspectives," *IEEE Transactions on Radiation and Plasma Medical Sciences*, vol. 8, no. 2, pp. 153–172, 2024.
- [21] P. Guo, C. Gong, X. Lin, Z. Yang, and Q. Zhang, "Exploring the adversarial frontier: Quantifying robustness via adversarial hypervolume," *IEEE Transactions on Emerging Topics in Computational Intelligence*, vol. 9, 2025.
- [22] H. Qiu, Y. Zeng, Q. Zheng, S. Guo, T. Zhang, and H. Li, "An efficient preprocessing-based approach to mitigate advanced adversarial attacks," *IEEE Transactions on Computers*, vol. 73, no. 3, pp. 645–655, 2021.
- [23] M. A. Ricci Lara, R. Echeveste, and E. Ferrante, "Addressing fairness in artificial intelligence for medical imaging," *nature communications*, vol. 13, no. 1, p. 4581, 2022.
- [24] H. Jang and F. Tong, "Improved modeling of human vision by incorporating robustness to blur in convolutional neural networks," *Nature Communications*, vol. 15, no. 1, p. 1989, 2024.
- [25] J. Kreis and J. Carroll, "Applications of adaptive optics imaging for studying conditions affecting the fovea," *Annual Review of Vision Science*, vol. 10, no. 1, pp. 239–262, 2024.
- [26] M. A. Shah, A. Kashaf, and B. Raj, "Training on foveated images improves robustness to adversarial attacks," in *Advances in Neural Information Processing Systems*, vol. 36, 2023, pp. 71 008–71 028.
- [27] O. Díaz, A. Rodríguez-Ruíz, and I. Sechopoulos, "Artificial intelligence for breast cancer detection: Technology, challenges, and prospects," *European Journal of Radiology*, vol. 175, p. 111457, 2024.
- [28] E. K. Park, S. Kwak, W. Lee, J. S. Choi, T. Kooi, and E.-K. Kim, "Impact of ai for digital breast tomosynthesis on breast cancer detection and interpretation time," *Radiology: Artificial Intelligence*, vol. 6, no. 3, p. e230318, 2024.
- [29] D. Resch, R. Lo Gullo, J. Teuwen, F. Semturs, J. Hummel, A. Resch, and K. Pinker, "Ai-enhanced mammography with digital breast tomosynthesis for breast cancer detection: Clinical value and comparison with human performance," *Radiology: Imaging Cancer*, vol. 6, no. 4, p. e230149, 2024.
- [30] J. L. Raya-Povedano, "Ai in breast cancer screening: A critical overview of what we know," *European Radiology*, vol. 34, no. 7, pp. 4774–4775, 2024.
- [31] W. S. Admass, Y. Y. Munaye, and A. O. Salau, "Integration of feature enhancement technique in google inception network for breast cancer detection and classification," *Journal of Big Data*, vol. 11, no. 1, p. 78, 2024.
- [32] M. R. Alom, F. Al Farid, M. A. Rahaman, A. Rahman, T. Debnath, A. S. M. Miah, and S. Mansor, "An explainable ai-driven deep neural network for accurate breast cancer detection from histopathological and ultrasound images," *Scientific Reports*, vol. 15, no. 1, p. 97718, 2025.
- [33] R. Sawyer-Lee, F. Gimenez, A. Hoogi, and D. Rubin, "Curated breast imaging subset of digital database for screening mammography (cbis-ddsm)," 2016. [Online]. Available: <https://www.cancerimagingarchive.net/collection/cbis-ddsm/>
- [34] Z. Cheng, F. Zhu, X.-Y. Zhang, and C.-L. Liu, "Adversarial training with distribution normalization and margin balance," *Pattern Recognition*, vol. 136, p. 109182, 2023.
- [35] M. G. Abdu-Aguye, M. Z. Zaheer, and K. Nandakumar, "Feature map purification for enhancing adversarial robustness of deep timeseries classifiers," in *Proceedings of the IEEE International Conference on Data Mining (ICDM)*, 2024, pp. 1–10.
- [36] H. Yu, P. Guo, and A. Sano, "Adawavenet: Adaptive wavelet network for time series analysis," *arXiv preprint arXiv:2405.11124*, 2024. [Online]. Available: <https://doi.org/10.48550/arXiv.2405.11124>
- [37] A. Zaytsev, M. Kovaleva, A. Natekin, E. Vorsin, V. Smirnov, G. Smirnov, O. Sidorshin, A. Senin, A. Dudin, and D. Berestnev, "Designing an attack-defense game: how to increase the robustness of financial transaction models via a competition," in *Proceedings of the IEEE International Conference on Data Mining (ICDM)*, 2024, pp. 600–609.
- [38] C. Cui, L. Li, H. Cai, Z. Fan, L. Zhang, T. Dan, J. Li, and J. Wang, "The chinese mammography database (cmmd): An online mammography database with biopsy confirmed types for machine diagnosis of breast," 2021. [Online]. Available: <https://doi.org/10.7937/tcia.eqde-4b16>
- [39] M. Buda, A. Saha, R. Walsh, S. Ghate, N. Li, A. Swiecicki, J. Y. Lo, J. Yang, and M. A. Mazurowski, "Breast cancer screening – digital breast tomosynthesis (bcs-dbt), version 5," 2020. [Online]. Available: <https://www.cancerimagingarchive.net/collection/breast-cancer-screening-dbt/>
- [40] I. Fursov, M. Morozov, N. Kaplounkaya, E. Kovtun, R. Rivera-Castro, G. Gusev, D. Babaev, I. Kireev, A. Zaytsev, and E. Burnaev, "Adversarial attacks on deep models for financial transaction records," in *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining (KDD '21)*. New York, NY, USA: Association for Computing Machinery, 2021, pp. 2868–2878.
- [41] W. Ji, S. Yu, J. Wu, K. Ma, C. Bian, Q. Bi, J. Li, H. Liu, L. Cheng, and Y. Zheng, "Learning calibrated medical image segmentation via multi-rater agreement modeling," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2021)*. Virtual: IEEE / Computer Vision Foundation, June 2021, pp. 12 341–12 351.
- [42] M. Z. Joel, S. Umrao, E. Chang, R. Choi, D. X. Yang, J. S. Duncan, A. Omuro, R. Herbst, H. M. Krumholz, and S. Aneja, "Using adversarial images to assess the robustness of deep learning models trained on diagnostic images in oncology," *JCO Clinical Cancer Informatics*, vol. 6, p. e2100170, 2022.
- [43] X. Chen and T. Luo, "Adversarial-robust transfer learning for medical imaging via domain assimilation," in *Advances in Knowledge Discovery and Data Mining — 28th Pacific-Asia Conference on Knowledge Discovery and Data Mining (PAKDD 2024), Proceedings*. Taipei, Taiwan: Springer, 2024, pp. 335–349.
- [44] D. Vasan and M. Hammoudeh, "Enhancing resilience against adversarial attacks in medical imaging using advanced feature transformation training," *Current Opinion in Biomedical Engineering*, vol. 32, p. 100561, 2024.
- [45] X. Li and D. Zhu, "Robust detection of adversarial attacks on medical images," in *2020 IEEE 17th International Symposium on Biomedical Imaging (ISBI)*. Iowa City, IA, USA: IEEE, 2020, pp. 1154–1158.